


# CAP5636 Artificial Intelligence

## Introduction




[These slides are based on slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. All materials available at <http://ai.berkeley.edu>.]




# Today

---

- What is artificial intelligence?
- What can AI do?
- What is this course?



# Sci-Fi AI?



# What is AI?

---

The science of making machines that:

# Rational Decisions

---


We'll use the term **rational** in a very specific, technical way:

- Rational: maximally achieving pre-defined goals
- Rationality only concerns what decisions are made  
(not the thought process behind them)
- Goals are expressed in terms of the **utility** of outcomes
- Being rational means **maximizing your expected utility**

A better title for this course would be:

**Computational Rationality**

# Maximize Your Expected Utility



# What About the Brain?


- Brains (human minds) are very good at making rational decisions, but not perfect
- Brains aren't as modular as software, so hard to reverse engineer!
- “Brains are to intelligence as wings are to flight”
- Lessons learned from the brain: memory and simulation are key to decision making







# Designing Rational Agents

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its (expected) **utility**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions
- **This course** is about:
  - General AI techniques for a variety of problem types
  - Learning to recognize when and how a new problem can be solved with an existing technique



# Pac-Man as an Agent



# Course Topics

- Part I: Making Decisions
  - Fast search / planning
  - Adversarial and uncertain search
- Part II: Reasoning under Uncertainty
  - Bayes' nets
  - Decision theory
  - Machine learning

