

Maximum Entropy Classifiers

Dr. Demetrios Glinos
University of Central Florida

CAP6640 – Computer Understanding of Natural Language

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Discriminative Probabilistic Models

- So far, we have examined "generative" models
 - language modeling
 - Naïve Bayes
 - HMM
- Increasing use of conditional or "discriminative" probabilistic models
 - in NLP, speech processing, IR (and machine learning generally)
 - give high accuracy performance
 - easy to incorporate lots of linguistically important features
 - support automatic building of language-independent NLP components

Generative Models

- Given some data $\{ (d,c) \}$ of paired observations d and hidden classes c
- Generative models compute *joint* probabilities over both the observed and the hidden variables

$$P(c, d)$$

- which are used to generate the observed data from the hidden
- examples: n-gram models, NB, HMM, PCFG

Discriminative Models


- Given some data $\{ (d,c) \}$ of paired observations d and hidden classes c
- Discriminative models compute *conditional* probabilities over the hidden variables given the data

$$P(c | d)$$


- which are used to generate the observed data from the hidden
- examples:
 - logistic regression, conditional loglinear or maximum entropy models, conditional random fields
 - also (although not directly probabilistic): perceptrons, NNs, SVMs

Bayes Nets

- Bayes net diagrams draw circles for random variables and lines for direct dependencies
- Some variables are observed; some are hidden
- Each node is a CPT over the incoming arcs
 - CPT serves as a small classifier



Generative



Discriminative

Joint vs. Conditional Likelihood

- **Joint model**

- computes joint probabilities $P(d, c)$
- tries to maximize joint likelihood
- trivial to choose weights: just use relative frequencies

- **Conditional model**

- computes conditional probabilities $P(c | d)$
- given the data, models only the conditional probability of the class
- tries to maximize the conditional likelihood
 - this is more difficult to do, as we shall see
 - but we do it to get increased accuracy

Usefulness of Discriminative Modeling

- Klein and Manning 2002, using Senseval-1 Data on a word sense disambiguation task:

Training Set	
Objective	Accuracy
Joint likelihood	86.8
Conditional likelihood	98.5

Test Set	
Objective	Accuracy
Joint likelihood	73.6
Conditional likelihood	76.1

- Comparison tests used the *same* word/class features and same smoothing

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Features


- Now, we *could* build a MEMM based on just the current word and previous tag

$$\hat{T} = \operatorname{argmax}_T \prod_i P(t_i | w_i, t_{i-1})$$

- But this would be no more accurate than the generative HMM model, and could even be worse
- The attraction of MEMMs is that we can use this formulation to incorporate additional, linguistically relevant features
 - The current word and previous tag are just 2 features
 - Many others features of the context are also possible


Incorporating More Features

- Typical features used in MEMMs
 - current word
 - neighboring words
 - previous tags
 - combinations of the above
 - plus additional features



Feature Templates

- Feature templates are used to specify combinations of features



- Example templates

$\langle t_i, w_{i-2} \rangle \langle t_i, w_{i-1} \rangle \langle t_i, w_i \rangle \langle t_i, w_{i+1} \rangle \langle t_i, w_{i+2} \rangle$

← *word-based*

$\langle t_i, t_{i-1} \rangle \langle t_i, t_{i-2}, t_{i-1} \rangle$

← *token-based*

$\langle t_i, t_{i-1}, w_i \rangle \langle t_i, w_{i-1}, w_i \rangle \langle t_i, w_i, w_{i+1} \rangle$

← *combined*

Example: Features generated

Given

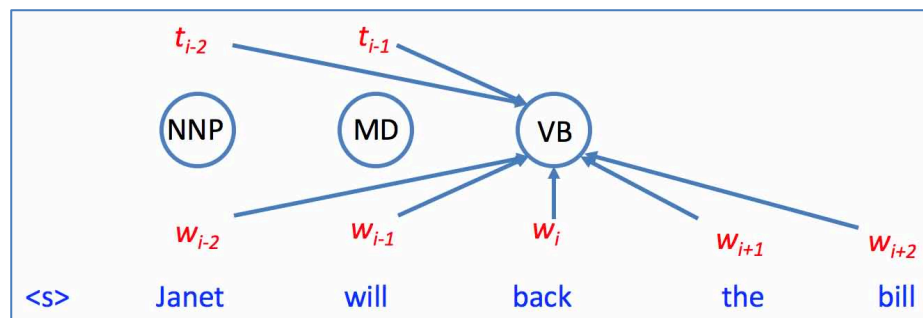
sentence "<s> Janet will back the bill"

Janet tagged NNP

will tagged MD

index i points to word "back"

Features generated



$t_i = \text{VB}$ and $w_{i-2} = \text{Janet}$

$t_i = \text{VB}$ and $w_{i-1} = \text{will}$

$t_i = \text{VB}$ and $w_i = \text{back}$

$t_i = \text{VB}$ and $w_{i+1} = \text{the}$

$t_i = \text{VB}$ and $w_{i+2} = \text{bill}$

$t_i = \text{VB}$ and $t_{i-1} = \text{MD}$

$t_i = \text{VB}$ and $t_{i-1} = \text{MD}$ and $t_{i-2} = \text{NNP}$

$t_i = \text{VB}$ and $w_{i-1} = \text{will}$ and $w_i = \text{back}$

$t_i = \text{VB}$ and $w_i = \text{back}$ and $w_{i+1} = \text{the}$

Word Signature Features

- MEMMs also typically use word signature features for unknown words
 - word-spelling properties
 - word shape
- Examples
 - w_i contains a particular prefix (from all prefixes of length ≤ 4)
 - w_i contains a particular suffix (from all suffixes of length ≤ 4)
 - w_i contains a number
 - w_i contains an upper case letter
 - w_i contains a hyphen
 - w_i is all upper case
 - w_i 's word shape
 - w_i 's short word shape
 - w_i is upper case and has a digit following a hyphen (e.g., CRC-12)
 - w_i is upper case and is followed within 3 words by Co., Inc., etc.

Word Shape Features

- Similar to regular expressions
- Basic word shape characteristics
 - map lower case letters to 'x', upper case to 'X', and digits to 'd'
 - examples
 - **X.X.X.** would match I.M.F.
 - **XXdd-dd** would match DC10-30
- Short word shape
 - same as above, but remove consecutive duplicate specification characters
 - examples
 - **Xd-d** would match DC10-30 and also B747-300

Feature Space

- MEMMs compute
 - Template features for every word seen in the training data set
 - Unknown word features for
 - all words in training set
 - or, just the low frequency ones (below some threshold)
- This produces a very large set of features
- **Feature cutoff**
 - features are not computed if they have a count < 5 in training set

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Feature-Based Linear Classifiers

- Linear classifiers
 - Linear function from feature sets $\{ f_i \}$ to classes $\{ c_j \}$
 - Assign a weight λ_i to each feature f_i
 - We consider each class for an observed datum d (feature set)
 - For the pair (c, d) , the features vote with their weights

$$vote(c_j) = \sum_i \lambda_i f_i(c_j, d)$$

- the winner is the class with the highest score

$$c^* = \operatorname{argmax}_{c_j} \sum_i \lambda_i f_i(c_j, d)$$

Methods for Choosing the Weights

- **Perceptron**
 - find a currently misclassified example, and nudge the weights in the direction of the correct classification
- **Margin-based methods**
 - e.g., Support Vector Machines
- **Exponential methods**
 - e.g., log-linear, **maxent**, logistic, Gibbs models
 - make a probabilistic model from the linear voting combination

$$P(c|d, \lambda) = \frac{\exp \sum_i \lambda_i f_i(c, d)}{\sum_{c'} \exp \sum_i \lambda_i f_i(c', d)}$$

← makes votes positive
← normalizes votes

- the **weights** are the **parameters** of the probability model

Example: Exponential Model

Given

vote(LOC | in Boulder) = 1.8

vote(PER | in Boulder) = - 0.6

vote(ORG | in Boulder) = 0.3

$$P(c|d, \lambda) = \frac{\exp \sum_i \lambda_i f_i(c, d)}{\sum_{c'} \exp \sum_i \lambda_i f_i(c', d)}$$

Then

$P(\text{LOC} \mid \text{in Boulder}, \lambda) = e^{1.8} / (e^{1.8} + e^{-0.6} + e^{0.3}) = 0.586$

$P(\text{PER} \mid \text{in Boulder}, \lambda) = e^{-0.6} / (e^{1.8} + e^{-0.6} + e^{0.3}) = 0.176$

$P(\text{ORG} \mid \text{in Boulder}, \lambda) = e^{0.3} / (e^{1.8} + e^{-0.6} + e^{0.3}) = 0.238$

→ relative ordering of vote results is preserved, but normalized to [0,1]

Note on Logistic Regression

- Exponential models
 - Goal of training is to choose the set of parameters $\{\lambda_i\}$ that **maximizes the conditional likelihood** of the data
 - To do this, we must construct not only classifications, but probability distributions over classifications
- Related to **logistic regression**
 - Maxent models in NLP are essentially the same as multiclass logistic regression models in statics or machine learning
 - parameterization is slightly different in a way that is useful for NLP-style models that have tons of sparse features
 - features are more general in that a feature is also a function of the class

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Features in NLP

- In NLP, a *feature* usually specifies
 1. an indicator function – a yes/no boolean matching function – of properties of the input
 2. and a particular class

$$f_i(c, d) \equiv [\Phi(d) \wedge c = c_j] \quad \leftarrow \text{value is 0 or 1}$$

- Each feature selects a data subset and suggests a label for it

Feature Expectations

- We will make use of two kinds of expectations
 - Empirical (actual) expectation of a feature

$$E(f_i, C) = \sum_{(c,d) \in \text{observed}(C,D)} f_i(c, d)$$

- Model (predicted) expectation of a feature

$$E(f_i, \lambda) = \sum_{(c,d) \in (C,D)} P(c, d) f_i(c, d)$$

- Goal of training a maxent model: to have the model expectations match the observed (empirical) expectations

Building a Maxent Model

- Features are often added during model development to target errors
 - Often, the easiest features to think of are the ones that indicate bad combinations
- Then, for any given feature weights, we wish to calculate
 - Data conditional likelihood
 - Derivative of the likelihood with respect to each feature weight
- We can then find the optimum feature weights (discussed later)

Exponential Model Likelihood

- Maximum (conditional) likelihood models
 - Given a model form, choose values of the parameters to maximize the (conditional) likelihood of the data

$$\log P(C|D, \lambda) = \sum_{(c,d) \in (C,D)} \log P(c|d, \lambda) = \sum_{(c,d) \in (C,D)} \log \frac{\exp \sum_i \lambda_i f_i(c, d)}{\sum_{c'} \exp \sum_i \lambda_i f_i(c', d)}$$

The Likelihood Value

- The (log) conditional likelihood of iid (independent, identically-distributed) data (C, D) according to a maxent model is a function of the data and the parameters λ

$$\log P(C|D, \lambda) = \log \prod_{(c,d) \in (C,D)} P(c|d, \lambda) = \sum_{(c,d) \in (C,D)} \log P(c|d, \lambda)$$

- If there aren't many values of c (i.e., data is sparse), it is easy to calculate

$$\log P(C|D, \lambda) = \sum_{(c,d) \in (C,D)} \log \frac{\exp \sum_i \lambda_i f_i(c, d)}{\sum_{c'} \exp \sum_i \lambda_i f_i(c', d)}$$

The Likelihood Value

- We can separate the last equation into two components

$$\begin{aligned} \log P(C|D, \lambda) = & \sum_{(c,d) \in (C,D)} \log \exp \sum_i \lambda_i f_i(c, d) \\ & - \sum_{(c,d) \in (C,D)} \log \sum_{c'} \exp \sum_i \lambda_i f_i(c', d) \end{aligned}$$

- Which is in the form

$$\log P(C|D, \lambda) = N(\lambda) - M(\lambda)$$

- The derivative is the difference between the derivatives of each component

Derivative (Part 1): Numerator

$$\begin{aligned}\frac{\partial N(\lambda)}{\partial \lambda_i} &= \frac{\partial}{\partial \lambda_i} \left(\sum_{(c,d) \in (C,D)} \log \exp \sum_i \lambda_i f_i(c,d) \right) \\ &= \frac{\partial}{\partial \lambda_i} \left(\sum_{(c,d) \in (C,D)} \sum_i \lambda_i f_i(c,d) \right) \\ &= \sum_{(c,d) \in (C,D)} \frac{\partial}{\partial \lambda_i} \sum_i \lambda_i f_i(c,d) \\ &= \sum_{(c,d) \in (C,D)} f_i(c,d)\end{aligned}$$

Thus, the derivative of the numerator is the empirical count(f_i, C)

Derivative (Part 2): Denominator

$$\begin{aligned}
 \frac{\partial M(\lambda)}{\partial \lambda_i} &= \frac{\partial}{\partial \lambda_i} \left(\sum_{(c,d) \in (C,D)} \log \sum_{c'} \exp \sum_i \lambda_i f_i(c', d) \right) \\
 &= \sum_{(c,d) \in (C,D)} \frac{1}{\sum_{c''} \exp \sum_i \lambda_i f_i(c'', d)} \frac{\partial}{\partial \lambda_i} \left(\sum_{c'} \exp \sum_i \lambda_i f_i(c', d) \right) \\
 &= \sum_{(c,d) \in (C,D)} \frac{1}{\sum_{c''} \exp \sum_i \lambda_i f_i(c'', d)} \sum_{c'} \exp \sum_i \lambda_i f_i(c', d) \frac{\partial}{\partial \lambda_i} \left(\sum_i \lambda_i f_i(c', d) \right) \\
 &= \sum_{(c,d) \in (C,D)} \sum_{c'} \frac{1}{\sum_{c''} \exp \sum_i \lambda_i f_i(c'', d)} \exp \sum_i \lambda_i f_i(c', d) \frac{\partial}{\partial \lambda_i} \left(\sum_i \lambda_i f_i(c', d) \right) \\
 &= \sum_{(c,d) \in (C,D)} \sum_{c'} P(c'|d, \lambda) f_i(c', d)
 \end{aligned}$$

Thus, the derivative of the denominator is the predicted count(f_i, λ)

Derivative (Part 3): Log Likelihood

- Combining the previous results, we have

$$\frac{\partial}{\partial \lambda_i} (\log P(C|D, \lambda)) = \text{actual count}(f_i, C) - \text{predicted count}(f_i, \lambda)$$

- The optimum parameters will be the ones for which each feature's **predicted expectation** equals its **empirical expectation**
- Optimal distribution
 - Always exists if feature counts are from actual data
 - Will always be unique (but parameters may not be unique)
- Such models are called maximum entropy models because the optimal parameters correspond to a model with maximum entropy
 - recall Shannon's definition of information entropy: $E(x) = -\ln \left(\frac{1}{P(x)} \right)$

Finding the Optimal Parameters


- We wish to choose parameters $\lambda_1, \lambda_2, \lambda_3, \dots$ that maximize the conditional log-likelihood of the training data

$$CLogLik(D) = \log P(C|D, \lambda)$$

- To be able to find the maximum, we have worked out how to calculate
 - the function value
 - and its partial derivatives (gradients) with respect to each model parameter

Finding the Optimal Parameters

- Use your favorite numerical optimization package
 - typically, one must minimize the negative of CLogLik



1. Gradient descent (GD); stochastic gradient descent (SGD)
2. Generalized iterative scaling (GIS) and improved iterative scaling (IIS)
3. Conjugate gradient (CG), maybe with preconditioning
4. Quasi-Newton methods: limited variable metric (LMVM) methods, e.g., L_BFGS

Most Likely Sequence

- Most likely sequence computed based on words within $\pm \ell$ words and the previous k tags

$$\hat{T} = \operatorname{argmax}_T \prod_i P(t_i | w_{i-\ell}^{i+\ell}, t_{i-k}^{i-1})$$

- Given the model (developed from a corpus)
 - Can use greedy algorithm
 - make a hard classification on the first word in the sentence, then on the second word, and so on
 - Can use the Viterbi algorithm
 - same approach as for HMM

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Bidirectionality

- Both HMM and MEMM work essentially left-to-right
 - even though Viterbi allows some influence on current tag by subsequent ones
- Conditional Random Fields (CRF)
 - allow explicit dependence on subsequent tags
 - more powerful
 - but at substantial computational cost
- Other approaches
 - Stanford tagger uses a bidirectional version of MEMM called a cyclic dependency network
 - can also turn any sequence model into a bidirectional model by performing multiple passes (e.g., first pass use only preceding tags, second pass, use tags on right as well)

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Morphologically Rich Languages

- Languages similar to English
 - HMM and MEMM tagger accuracies comparable to English
- Morphologically rich languages
 - e.g., Czech, Hungarian, Turkish
 - for comparable size dictionary
 - Hungarian contains 2x word types as English
 - Turkish contains 4x word types as English
 - larger vocabularies mean more unknown word types and result in degraded performance
 - word morphology also encodes more information, like case and gender, so POS taggers need to label these features as well
 - tagsets can be 4 to 10 times larger than the 50-100 we use for English

Non-Segmented Languages

- Languages like Chinese do not segment written words
 - word segmentation generally applied before tagging
 - but some methods perform segmentation and tagging jointly
 - unknown words are a significant problem
 - despite short (~2.4 characters/unknown ~7.7 char/unk for English)
 - in English, unknown words tend to be proper nouns
 - in Chinese, unknowns tend to be common nouns and verbs
 - features for unknowns
 - use prefixes and suffixes (same as in English)
 - but also use radicals (not used in English)

Today

- Discriminative Models
- Features for Discriminative Models
- Feature-Based Linear Classifiers
- Building a Maxent Model
- Bidirectionality
- POS Tagging for Other Languages
- POS Tagging Performance

Baseline Tagging Performance

- Performance metric: Tag accuracy
 - how many tags are correct
 - about 97% currently
- Baseline method
 - tag every word with its most frequent tag
 - tag unknown words as common nouns
- Baseline performance is 90%
 - reason is many words are unambiguous
 - unambiguous words (like "the", "a", etc.) and for punctuation marks are included in the accuracy calculation

Comparison of POS Tagging Methods

- Rough accuracies for various POS tagging methods
 - on all words
 - on unknown words

Method	Accuracy (total)	Accuracy (unknowns)
Most frequent tag	~90%	~50%
Trigram HMM	~95%	~55%
Maxent $P(t w)$	93.7%	82.6%
MEMM tagger	96.9%	86.9%
Bidirectional dependencies	97.2%	90%
Upper bound	~98% (human agreement)	