

Advanced Dialog Systems

Dr. Demetrios Glinos
University of Central Florida

CAP6640 – Computer Understanding of Natural Language

Today

- Dialog-State Systems
- Dialog Acts
- Interpreting Dialoge Acts
- Dialog Policy
- Natural Language Generation

Excerpt: Abbott & Costello's "Who's on First"

C : I want you to tell me the names of the fellows on the St. Louis team.

A : I'm telling you. Who's on first, What's on second, I Don't Know is on third--

C : You know the fellows' names?

A : Yes.

C : Well, then who's playing first?

A : Yes.

C : I mean the fellow's name on first base.

A : Who.

C : The fellow playin' first base.

A : Who.

C : The guy on first base.

A : Who is on first.

C : Well, what are you askin' me for?

A : I'm not asking you--I'm telling you. Who is on first.

C : I'm asking you--who's on first?

A : That's the man's name.

C : That's who's name?

A : Yes.

Speech or Dialog Acts

- **Speech or dialog acts**
 - what we do in conversation
 - making a statement, asking a question, giving an order, uttering an exclamation
- GUS-style frame-based dialog systems cannot handle different dialog acts
 - but frames and slots will be useful here, too
- **Dialog-state systems**
 - also called belief-state or information-state systems
 - also based on filling frames and slots
 - but can understand and generate a variety of dialog acts
 - can track who has the initiative
 - what to say next is based on a **dialog policy** (more complex than filling the next slot)

Dialog-State System Architecture

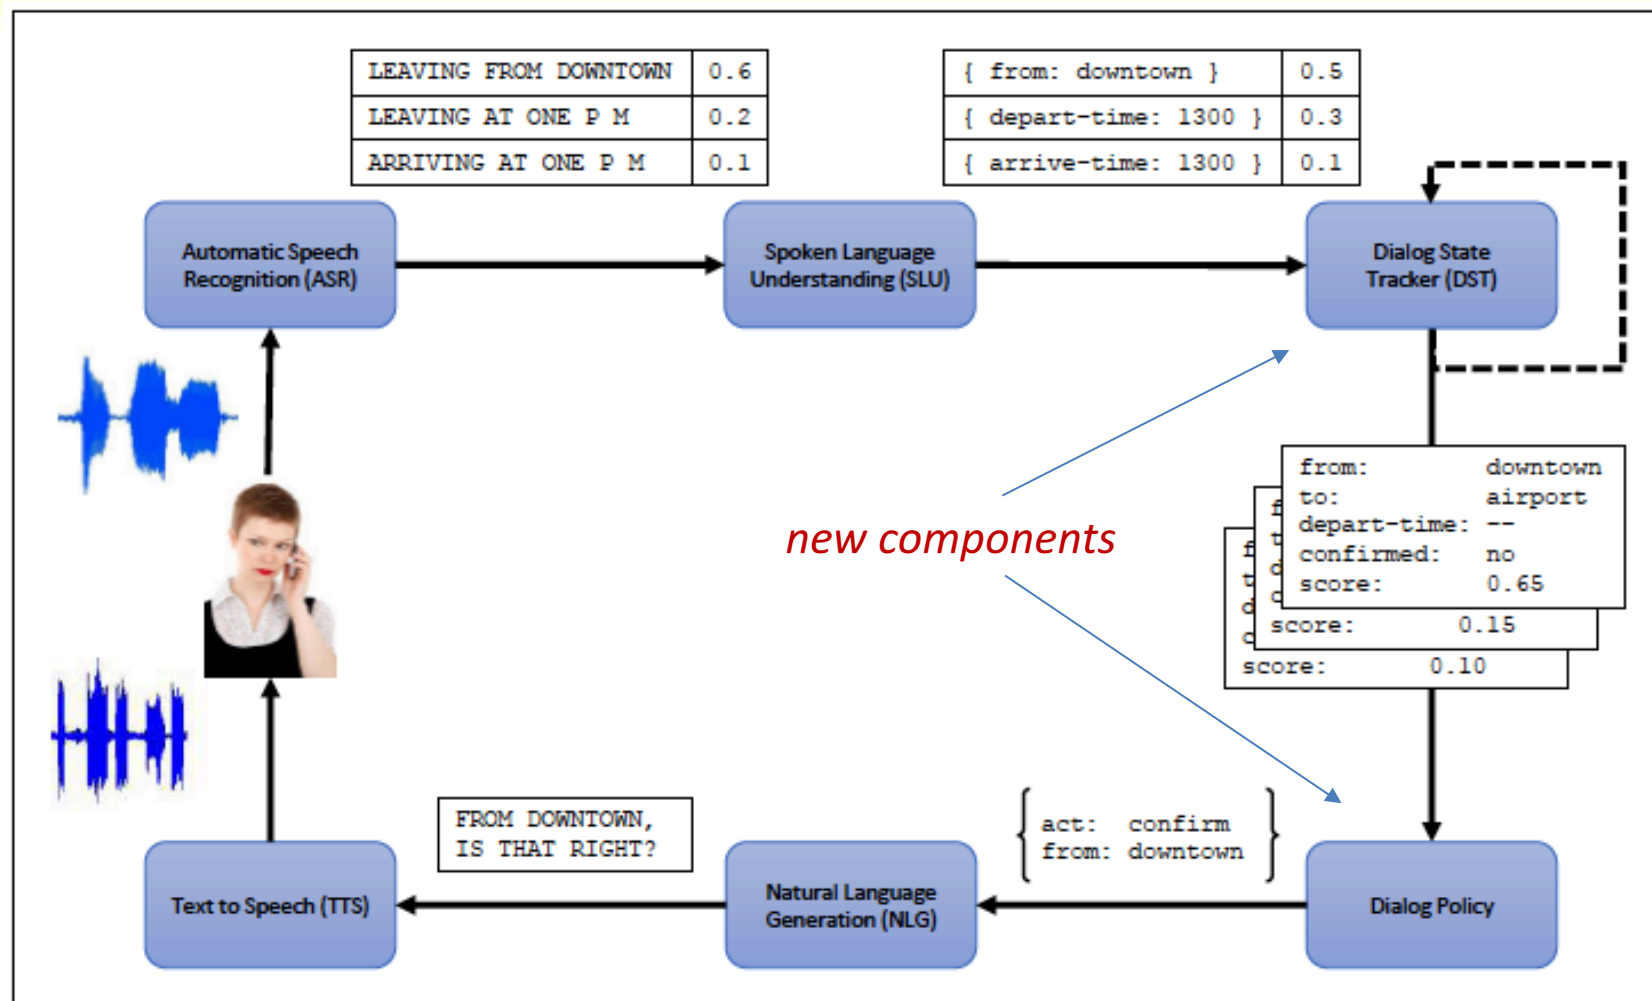


Figure 30.1 Architecture of a dialog-state system for task-oriented dialog from Williams et al. (2016).

source: J&M (3d Ed. draft)

Today

- Dialog-State Systems
- **Dialog Acts**
- Interpreting Dialoge Acts
- Dialog Policy
- Natural Language Generation

Dialog Acts

- Key insight: Utterances are actions, commonly called speech acts
- **Speech acts** (Bach and Harnish, 1979):
 - Constatives:** committing the speaker to something being the case
e.g., *answering, claiming, confirming, denying, denying, stating*
 - Directives:** attempts by the speaker to get the addressee to do something
e.g., *ordering, asking, advising, forbidding, inviting, requesting*
 - Commissives:** committing the speaker to some future course of action
e.g., *promising, planning, vowing, betting, opposing*
 - Acknowledgments:** expressing the speaker's attitude regarding the hearer with respect to some social action
e.g., *apologizing, greeting, thanking, accepting an acknowledgment*

Common Ground

- Dialog is not a series of independent speech acts
- It is a collective act performed by speaker and hearer
- **Common ground**
 - the set of things that are mutually believed by both speakers
 - must be continually established
 - hearer must *ground* the speaker's utterances
 - i.e., acknowledge them to make it clear that the hearer has understood the speaker's meaning and intention
- **Principle of closure:** Agents performing an action require evidence, sufficient for current purposes, that they have succeeded in performing it
 - humans need this, too
 - this is why a well-designed elevator button lights up when we press it
 - grounding is also important to tell the speaker that he has *not* succeeded

Grounding Methods

- A continuum of grounding methods, ordered from weakest to strongest (Clark and Schaefer, 1989):

Continued attention	B shows she is continuing to attend and therefore remains satisfied with A's presentation
Next contribution	B starts in on the next relevant contribution
Acknowledgment	B nods or says a continuer like <i>uh-huh</i> , <i>yeah</i> , or the like, or an assessment like <i>that's great</i>
Demonstration	B demonstrates all or part of what she has understood A to mean, by reformulating (paraphrasing) A's utterance or by collaborative completion of A's utterance
Display	B displays verbatim all or part of A's presentation

Examples of Grounding

C₁: ...I need to travel in May.
 A₁: And, what day in May did you want to travel? ← *display*
 C₂: OK uh I need to be there for a meeting that's from the 12th to the 15th.
 A₂: And you're flying into what city?
 C₃: Seattle. ← *next contribution*
 A₃: And what time would you like to leave Pittsburgh? ← *acknowledgment*
 C₄: Uh hmm I don't think there's many options for non-stop.
 A₄: Right. There's three non-stops today.
 C₅: What are they?
 A₅: The first one departs PGH at 10:00am arrives Seattle at 12:05 their time. The second flight departs PGH at 5:55pm, arrives Seattle at 8pm. And the last flight departs PGH at 8:15pm arrives Seattle at 10:28pm.
 C₆: OK I'll take the 5ish flight on the night before on the 11th.
 A₆: On the 11th? OK. Departing at 5:55pm arrives Seattle at 8pm, U.S. Air flight 115.
 C₇: OK.

Figure 30.2 Part of a conversation between a travel agent (A) and client (C).

source: J&M (3d Ed. draft)

Speech Acts in Practical Systems

- Practical systems must distinguish
 - statements
 - directives asking for information
 - directives that are orders to do something
- Speech acts are important here, too
- Example

Without acknowledgment	With acknowledgement
System: Do you wish to review more of your personal profile?	System: Do you wish to review more of your personal profile?
Caller: No.	Caller: No.
System: What's next?	System: Okay, what's next?

Dialog Acts

- **Dialog acts**
 - combine speech acts and grounding
 - serve as tags for entire sentences
 - tagset is generally custom built for each particular application

Tag	Sys	User	Description	} slot fillers
HELLO($a = x, b = y, \dots$)	✓	✓	Open a dialog and give info $a = x, b = y, \dots$	
INFORM($a = x, b = y, \dots$)	✓	✓	Give info $a = x, b = y, \dots$	
REQUEST($a, b = x, \dots$)	✓	✓	Request value for a given $b = x, \dots$	
REQALTS($a = x, \dots$)	✗	✓	Request alternative with $a = x, \dots$	
CONFIRM($a = x, b = y, \dots$)	✓	✓	Explicitly confirm $a = x, b = y, \dots$	
CONFREQ($a = x, \dots, d$)	✓	✗	Implicitly confirm $a = x, \dots$ and request value of d	
SELECT($a = x, a = y$)	✓	✗	Implicitly confirm $a = x, \dots$ and request value of d	
AFFIRM($a = x, b = y, \dots$)	✓	✓	Affirm and give further info $a = x, b = y, \dots$	
NEGATE($a = x$)	✗	✓	Negate and give corrected value $a = x$	
DENY($a = x$)	✗	✓	Deny that $a = x$	
BYE()	✓	✓	Close a dialog	

Figure 30.4 Dialogue acts used by the HIS restaurant recommendation system of [Young et al. \(2010\)](#). The Sys and User columns indicate which acts are valid as system outputs and user inputs, respectively.

source: J&M (3d Ed. draft)

Example: Sample HIS System Dialog

Utterance	Dialogue act
U: Hi, I am looking for somewhere to eat.	<code>hello(task = find,type=restaurant)</code>
S: You are looking for a restaurant. What type of food do you like?	<code>confreq(type = restaurant, food)</code>
U: I'd like an Italian somewhere near the museum.	<code>inform(food = Italian, near=museum)</code>
S: Roma is a nice Italian restaurant near the museum.	<code>inform(name = "Roma", type = restaurant, food = Italian, near = museum)</code>
U: Is it reasonably priced?	<code>confirm(pricerange = moderate)</code>
S: Yes, Roma is in the moderate price range.	<code>affirm(name = "Roma", pricerange = moderate)</code>
U: What is the phone number?	<code>request(phone)</code>
S: The number of Roma is 385456.	<code>inform(name = "Roma", phone = "385456")</code>
U: Ok, thank you goodbye.	<code>bye()</code>

Figure 30.5 A sample dialog from the HIS System of [Young et al. \(2010\)](#) using the dialog acts in Fig. 30.4.

source: J&M (3d Ed. draft)

Conversational Analysis

- **Conversational analysis (CA)**
 - analyzes structure of conversations by examining adjacency pairs
 - consists of a "first pair part" and a "second pair part"
 - because dialog acts are not independent
 - examples:
 - question and answer
 - proposal and acceptance (or rejection)
 - compliment and downplayer
 - greeting and greeting
- adjacency pair structure can help model decide which action to take

Adjacency Pair Separation

- adjacency pair components can be separated by a side sequence or a pre-sequence

- side sequence

User: What do you have going to UNKNOWN_WORD on the 5th?

System: Let's see, going *where* on the 5th?

User: Going to Hong Kong.

System: OK, here are some flights ...

- pre-sequence

User: Can you make train reservations?

System: Yes I can.

User: Great, I'd like to reserve a seat on the 4pm train to New York.

Today

- Dialog-State Systems
- Dialog Acts
- Interpreting Dialoge Acts
- Dialog Policy
- Natural Language Generation

Interpreting Dialogue Acts

- **Dialog-state tracker (DST)** determines
 - current fillers for the current frame
 - current set of user constraints
 - the user's most recent dialog act
- Example (Mrksic et al., 2017)

Conversation	DST output
User: I'm looking for a cheaper restaurant	inform(price=cheap)
System: Sure. What kind – and where?	
User: Thai food, somewhere downtown	inform(price=cheap, food=Thai, area=centre)
System: The House serves cheap Thai food	
User: Where is it?	inform(price=cheap, food=Thai, area=centre); request(address)
System: The House is at 106 Regent Street	

Using Surface Syntax

- Surface syntax seems informative
 - yes-no question: auxiliary verb precedes subject
e.g., *Will breakfast be served on the flight?*
 - statement: no aux-inversion
e.g., *I don't care about lunch.*
 - command: no syntactic subject
e.g., *Show me flights from Orlando to New York.*
- But there are common exceptions
 - request disguised as a question: *Can you give me flights from Miami to Austin?*
 - question disguised as a statement: *You said you wanted to travel on Friday?*

Learning Dialog Actions

- Dialog actions constrain the possible slots and values involved in an utterance

I'd like Cantonese food near Soho

→ `inform(food=cantonese, area=soho)`

- First-pass classifier chooses dialog act (*inform* above)
 - standard supervised machine learning on labeled corpus
 - lots of features, including unigrams, bigrams, etc.
 - e.g. show me → request, are there → question
- Second-pass classifier fills appropriate slots
 - standard supervised machine learning on labeled corpus
 - e.g., CRF or RNN-based IOB tagging
 - or a multinomial classifier for all possible slot-value pairs (since ontology finite)

Detecting Correction Acts

- **User correction acts**
 - When a system fails to recognize or misunderstands an utterance, the user generally repeats or reformulates the utterance
- If system contains an ASR component, detection is often more difficult
 - users typically use *hyperarticulation* (exaggerated energy, duration, etc.) when stating corrections
 - e.g., *I said BAL-TI-MORE, not Boston*
- Corrective utterances are typically
 - exact repetitions
 - repetitions with one or more words omitted
 - paraphrases
 - these can be learned by a standard machine learning classifier

Today

- Dialog-State Systems
- Dialog Acts
- Interpreting Dialoge Acts
- **Dialog Policy**
- Natural Language Generation

Dialog Policy

- **Dialog policy**
 - determines the next action the system should take
 - i.e., what dialog act the system should generate
- Running example: Confirmation and Rejection
 - systems make mistakes and often need to be sure they have correctly understood the user
 - methods
 - **confirming** understandings with the user
 - **rejecting** utterances that the system is likely to have misunderstood

Confirmation Methods

- **explicit confirmation** (generally sounds non-natural and non-human)

S: Which city do you want to leave from?
U: Baltimore.
S: **Do you want to leave from Baltimore?**
U: Yes.
U: I'd like to fly from Denver Colorado to New York City on September twenty first in the morning on United Airlines
S: **Let's see then. I have you going from Denver Colorado to New York on September twenty first. Is that correct?**
U: Yes

- **implicit confirmation** (generally sounds much more conversationally natural)

demonstration →

U: I want to travel to Berlin
S: **When do you want to travel to Berlin?**

display →

U2: Hi I'd like to fly to Seattle Tuesday Morning
A3: **Traveling to Seattle on Tuesday, August eleventh in the morning.**
Your full name?

source: J&M (3d Ed. draft), Ch. 30

Rejection Strategies

- **simple rejection:** *I'm sorry, I didn't understand that.*
- **progressive prompting (escalating detail):**
 - instead of repeating the same question

System: When would you like to leave?
Caller: Well, um, I need to be in New York in time for the first World Series game.
System: <reject>. Sorry, I didn't get that. Please say the month and day you'd like to leave.
Caller: I wanna go on October fifteenth.

source: J&M (3d Ed. draft), Ch. 30

- **rapid reprompting:** *I'm sorry?* or *What was that?*
 - used for the first rejection instead of a detailed follow-up question
 - preferred by users

Features for Decision Policy

- Decision (in our running example) is to choose:
 - explicit confirmation
 - implicit confirmation
 - rejection
- Common features
 - ASR confidence in what is recognized from utterance
 - cost of making an error (e.g., when booking flight or moving money)
- Example: using thresholds α , β , and γ

$< \alpha$	low confidence	reject
$\geq \alpha$ and $< \beta$	some confidence	confirm explicitly
$\geq \beta$ and $< \gamma$	high confidence	confirm implicitly
$\geq \gamma$	very high confidence	do not confirm at all

A Simple Policy Based on Local Context

- Basic policy:
 - choose the best system action A_i based on the entire dialog state
 - can include entire history of conversation by both parties

$$\widehat{A}_i = \operatorname{argmax}_{A_i \in A} P(A_i \mid A_1, U_1, \dots, A_{i-1}, U_{i-1})$$

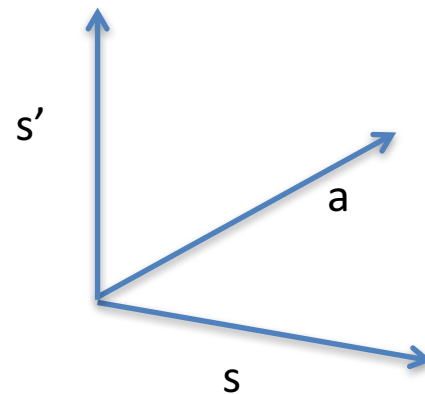
- Now, there are many different conversational paths that could lead to the same set of filled slots, so we can simplify by using the state of the affected frame

$$\widehat{A}_i = \operatorname{argmax}_{A_i \in A} P(A_i \mid \text{Frame}_{i-1}, A_{i-1}, U_{i-1})$$

- Given a sufficiently large corpus, we can estimate these probabilities using a classifier
- In practice, often need to build simulators to generate artificial conversations on which to train

Markov Decision Processes

- **Markov Decision Process (MDP)**
 - A set of **states** $s \in S$
 - A set of **actions** $a \in A$
 - A **transition model** $P(s' | s, a)$
 - A **reward function** $R(s)$
 - A **start state** s_0
 - Possibly one or more **terminal states**
- Transitions have the **Markov property**
 - $P(s' | s, a)$ does not depend on how the agent got to state s
 - we can visualize this a 3-D table of probabilities
- A solution to an MDP is a **policy, π**
 - $\pi(s)$ is the action to take for state s
- Executing a policy produces nondeterministic (stochastic) results
- An **optimal policy** must *maximize expected utility*



Modeling Dialog as an MDP

- States
 - can include entire history, but state space would be exceptionally large
 - so, typically use
 - same features as in our simplified model (frame + last turn)
 - plus ASR confidence, etc
- Actions
 - the particular speech acts to generate
 - also possibly performing a database query to find information (e.g., for flights)
- Reward function: $R = -(w_i n_i + w_e n_e + w_f n_f)$

$$\begin{aligned}\text{where } n_i &= \#turns \\ n_e &= \#errors \\ n_f &= \#slots\ filled \\ w_i &= weights\end{aligned}$$

Learning the Optimal Policy

We can use **reinforcement learning**

Representing a dialog as a trajectory in state space $s_1 \xrightarrow{a_1, r_1} s_2 \xrightarrow{a_2, r_2} s_3 \xrightarrow{a_3, r_3} s_4 \dots$

The cumulative reward for a given state sequence, using discount factor $\gamma \in [0,1)$

$$Q([s_0, a_0, s_1, a_1, s_2, a_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

The expected value of the cumulative reward is given by the Bellman equation

$$Q(s, a) = R(s) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q(s', a')$$

i.e., expected cumulative reward for a given state/action pair is

- the immediate reward for the state
- plus expected discounted utility of all possible next states s'
- weighted by the probability of transitioning to such state s'

Solving the Model

- To solve the Bellman equation for the dialog action, we need to know
 - $P(s'|s, a)$
 - $R(s)$
- We can compute maximum likelihood estimates for these values from a labeled training corpus
- Given these estimates, we can solve the model using **value iteration**
 - start with zero estimates for values
 - iteratively compute revised values until convergence
- **Optimal policy** is to move in the direction of the highest next state value
 - by taking the action with the highest probability of reaching that next state

Problems with MDPs for Dialog

- MDP only practical for small models
- Can extend using Partially Observable MDPs (POMDPs)
 - add extra latent variables to represent system's uncertainty about the true state of the dialog
- Both MDPs and POMDPs suffer from computational complexity due to their reliance on simulations that don't reflect true user behavior
- Recent efforts focus on applying reinforcement learning to deep neural networks, but there is no consensus yet on a standard paradigm

Today

- Dialog-State Systems
- Dialog Acts
- Interpreting Dialoge Acts
- Dialog Policy
- Natural Language Generation

NLG Component

- Natural language generation (NLG)
 - performed once a dialog action has been decided
 - modeled in two stages
 - content planning
 - we can assume that this is determined by the policy, when it chooses the action
 - and maybe also some additional attributes (slots and values) that the system wishes to implicitly confirm to the user
 - sentence realization
 - determining how to express the content in natural language

Example: Sentence Realization

```

{
  act query
  content depart_time
  depart_date {
    year 2000
    month 10
    day 5
  }
  depart_airport BOS
}

```

← *dialog act*
 ← *argument to dialog act*
 } *additional slot information to include*

=> What time on October fifth would you like to leave Boston?

Figure 30.6 An input frame to NLG and a resulting output sentence, in the Communicator system of Oh and Rudnicky (2000).

source: J&M (3d Ed. draft)

Steps in Sentence Realization

1. Generate **delexicalized** string

- basically, a generic sentence form, with variables for specific words
 - typically for named entities, dates, times, etc.
 - e.g., *What time on [departure date] would you like to leave [depart_airport] ?*

2. **Relexicalize**

- fill in variables from the frame provided by the content planner

Learning Delexicalized Strings

- Use human-human dialogs labeled with dialog acts and slots
- Delexicalize entire corpus
- Divide into separate corpora, one for each dialog act
- Example training sentences:

*And what time would you like to leave depart_city ?
When would you like to leave depart_city ?
When would you like to leave ?
What time do you want to leave on depart_date ?
OK, on depart_date, what time do you want to leave ?*

- Train a distinct N-gram or neural language model for each dialog act.
- **Usage:** Given dialog act QUERY DEPART_TIME, system will
 - sample random sentences from this model
 - rate them based on whether too short, too long, repeat or lacking slots, etc.
 - choose the highest scoring realization

Targeted Clarification

- Clarification questions
 - e.g., when ASR fails to understand some part of user's utterance
- Humans typically use targeted clarification questions that repeat elements of the misunderstanding

User: What do you have going to UNKNOWN_WORD on the 5th ?
System: Going where on the 5th ?

- Typically better received than generic rejection responses like "Please repeat."
- Creating targeted clarification questions
 - Can hand-craft rules: e.g., "going to UNKNOWN_WORD" → "going where"
 - Can also build a classifier to guess which slots might have been misrecognized